

# Chapter 2 Markov Decision Processes

COMP 3270 Artificial Intelligence

Dirk Schnieders

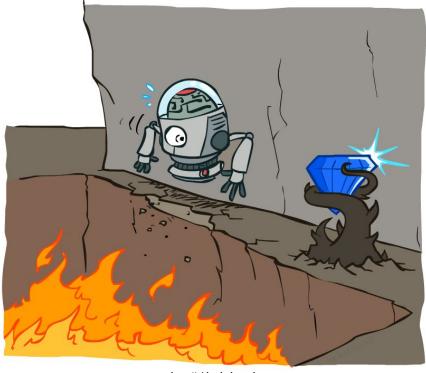
### Sequential Decision Problem

- In a sequential decision problem the agent's utility depends on a sequence of decisions
- Sequential decision problems incorporate utilities, uncertainty and sensing
- Optimal behavior balances the risks and rewards of acting in an uncertain environment



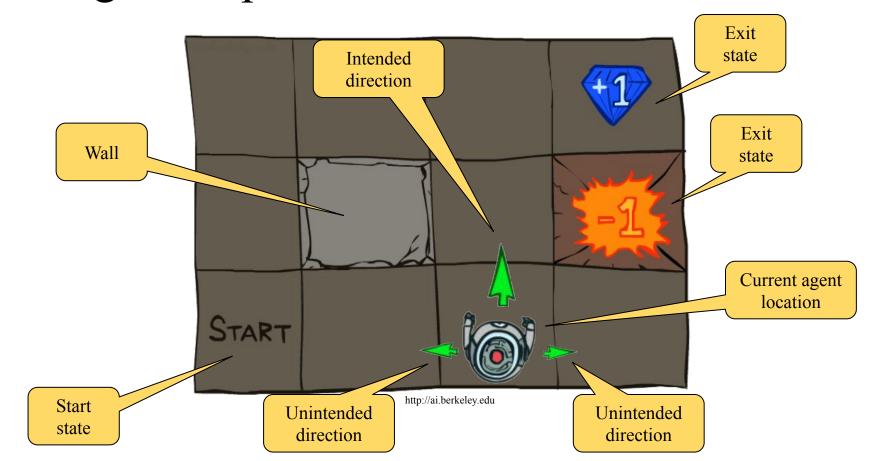
COMP 3270

#### Non-Deterministic Search



http://ai.berkeley.edu

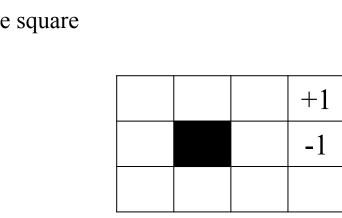
#### Running Example: Grid World



# Grid World

**COMP 3270** 

- Maze-like problem
  - Agent lives in a grid
  - Walls block the agent's path
- Unreliable actions
  - Each action achieves the intended effect 80% of the time
  - o 10% of the time the action moves the agent at right angles (i.e., 90°) to the intended direction
  - If the agent bumps into a wall it stays in the same square
- The agent receives rewards for each action
  - Small "living" rewards (can be negative)
  - Big rewards come at the end (good or bad)
- Goal states
  - o +1 and -1 are goal states
  - Only action available: exit action
- Goal: maximize sum of rewards



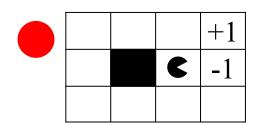
Example of stochastic motion:

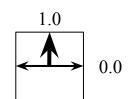
Intended direction is north

0.8

0.1

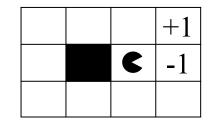
#### Deterministic vs. stochastic motion





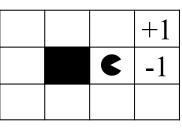
	•	+1
		-1

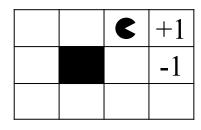


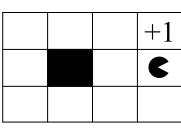


0.8

0.1





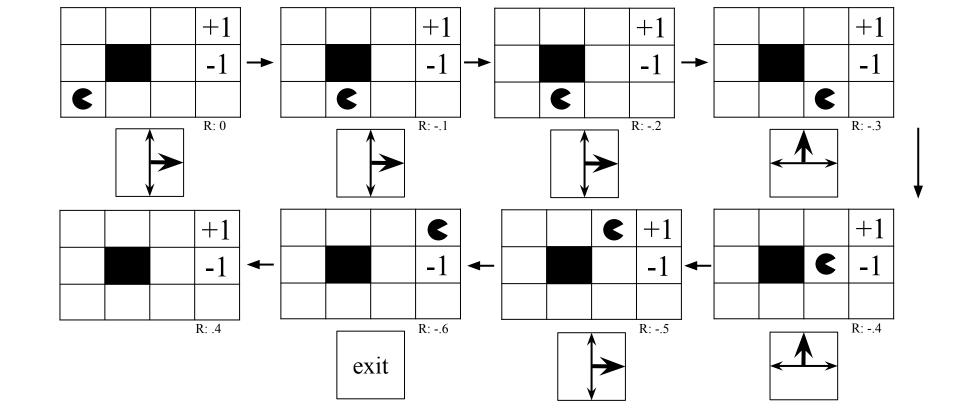


0.1

0.8

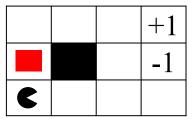
0.1

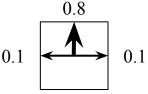
# Example episode in grid world



#### Quiz - North North

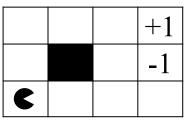
- Consider the fixed action sequence [North, North]
  - What is the probability of reaching the highlighted state from the start state with this action sequence?





#### Quiz - East East

• Which squares can be reached from the start state by the action sequence [East, East] and with what probabilities?



#### Transition model

- The transition model T(s,a,s') describes the outcome of each action in each state
- The outcome is stochastic and we write P(s'|s,a) to denote the probability of reaching state s' if action a is done in state s
- We assume that transitions are Markovian
  - I.e., the probability of reaching s' from s depends only on s and not on the history of earlier states
  - The Markov property is named after the Russian mathematician Andrey Markov



Andrey Markov (1856-1922)

#### Markov Decision Processes

- A sequential decision problem for a fully observable, stochastic environment with a Markovian transition model and additive rewards is called a Markov decision process (MDP)
- It is defined by
  - $\circ \quad \text{A set of states } s \in S$ 
    - A set of actions  $a \in A$
    - A transition function T(s,a,s')
    - Probability that a from s leads to s', i.e., P(s'|s,a)
    - A reward function R(s,a,s')
    - $\circ$  A start state  $s_0$
    - A terminal state (optional)

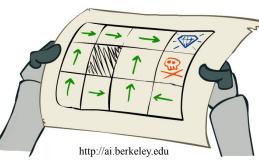
#### How to solve MDPs?

A fixed action sequence?

# Policy

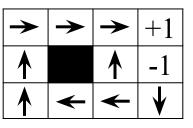
http://ai.berkeley.edu

- How to solve MDPs?
  - A fixed action sequence won't solve the problem because the agent might end up in a state other than the goal
  - A solution must specify what the agent should do for any state that the agent might reach
  - $\circ$  A solution of this kind is called a policy  $\pi$ 
    - $\pi(s)$  is the action recommended by the policy  $\pi$  for the state s

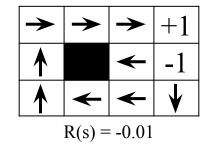


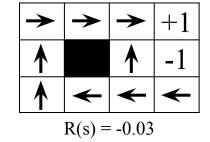
# Optimal Policy

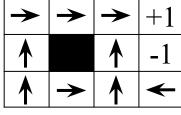
- Each time a given policy is executed, the stochastic nature of the environment may lead to a different environment history
- The quality of a policy is therefore measured by the expected utility of the possible environment histories generated by that policy
- An optimal policy  $\pi^*$  is a policy that yields the highest expected utility
- Example:  $\pi^*$  for Grid World with R(s) = -0.04 ( $\forall$  nonterminal s)



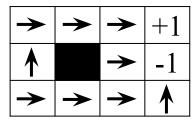
# Optimal Policy - Example







$$R(s) = -0.4$$



R(s) = -2.0

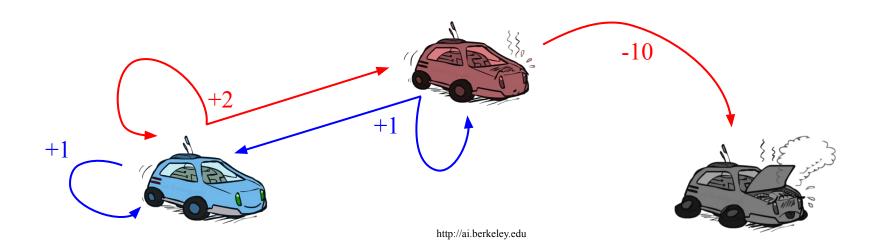
# Example: Racing Car



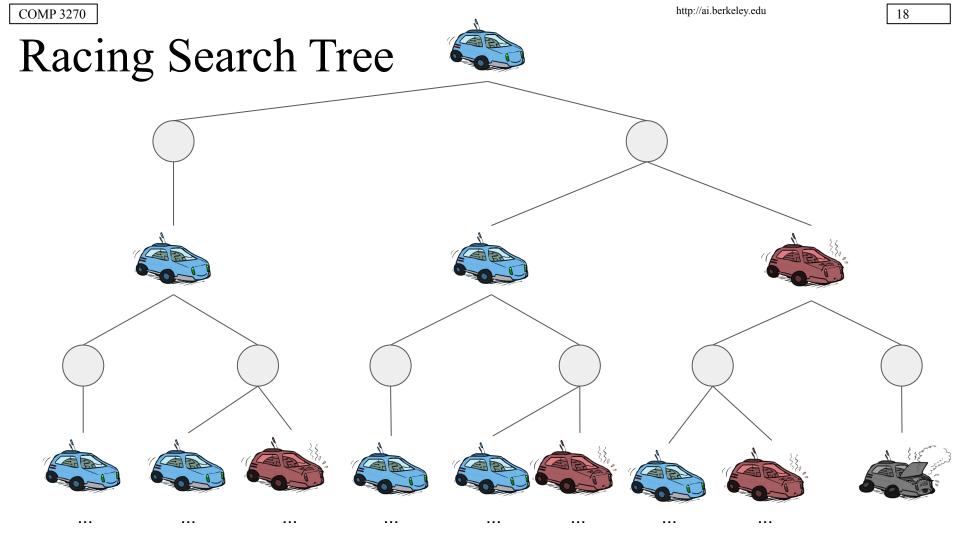
http://ai.berkeley.edu

## Example: Racing Car

- A robot car wants to travel far, quickly
- Three states: Cool, Warm, Overheated
- Two actions: Slow (+1), Fast (+2)

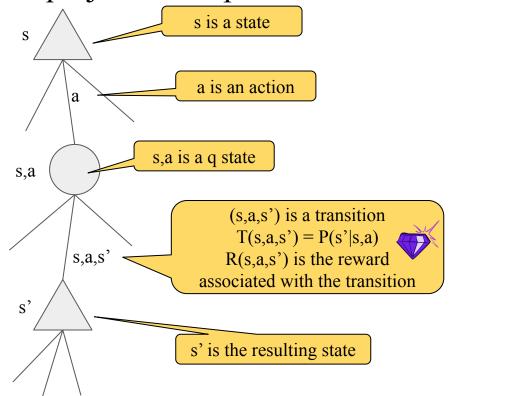






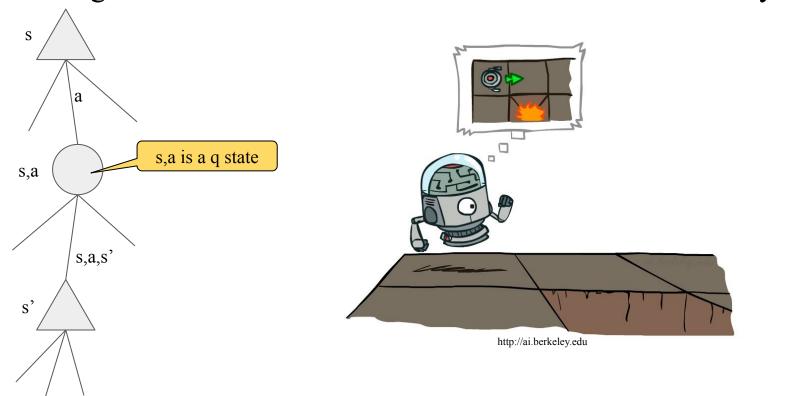
#### MDP Search Tree

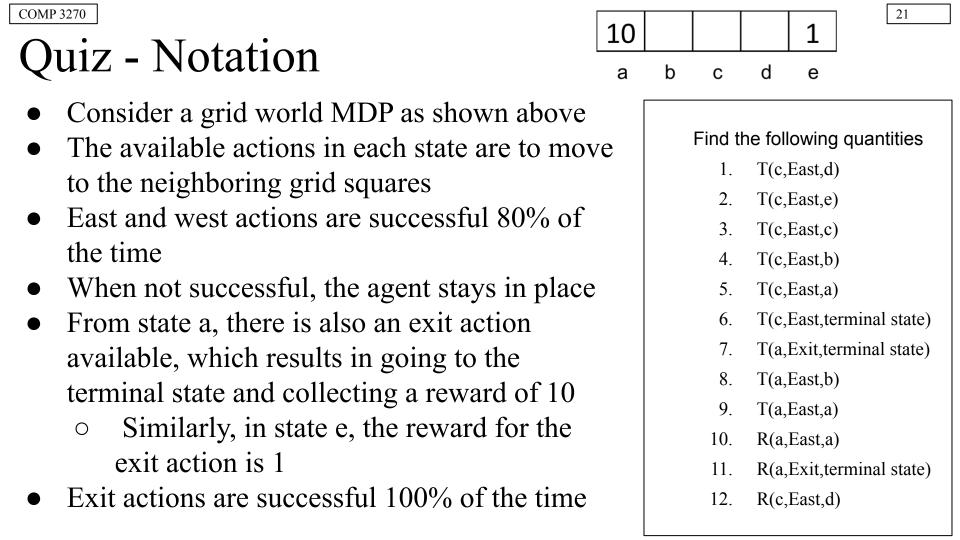
• Each MDP state projects an expectimax-like search tree



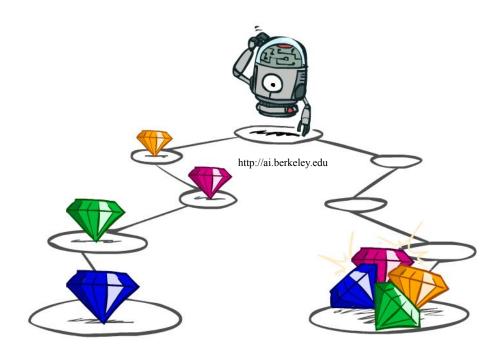
#### Q state

• The agent has committed to the action but has not done it yet



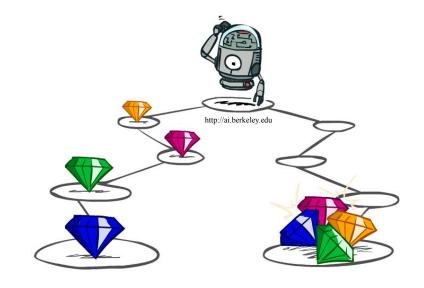


# Utility of State Sequences



#### Utility of State Sequences

- How to calculate the utility of state sequences?
  - What preferences should an agent have over reward sequences?
- More or less?
  - $\circ$  [1, 2, 2] or [2, 3, 4]
- Now or later?
  - o [0, 0, 1] or [1, 0, 0]



### Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution
  - Values of rewards decay exponentially

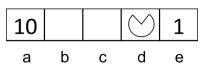


# Quiz - [1,2,3] vs. [3,2,1]

- Consider the sequences [1,2,3] and [3,2,1]
- Let  $\gamma = 0.5$ , which sequence has a higher utility?

#### Quiz - Optimal Action 1

- Consider the same grid world MDP as in one of the previous quiz
- Let actions always be successful
- Let the discount factor be  $\gamma = 0.1$ 
  - Q1: What is the optimal action in the state d?
- Now let the discount factor be  $\gamma = 0.9999$ 
  - Q2: What is the optimal action in the state d?



## Quiz - Optimal Action 2

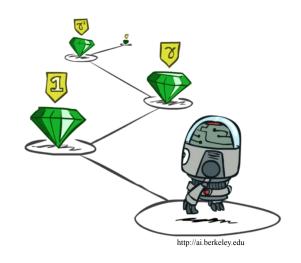
• Consider the following 101 x 3 world

+50	-1	-1	-1	• • •	-1	-1	-1	-1
Start				• • •				
-50	+1	+1	+1		+1	+1	+1	+1

- In the start state, the agent has a choice of two deterministic actions up or down but in the other states the agent has one deterministic action
- Assuming a discounted reward function, for what values of the discount  $\gamma$  should the agent choose up and for which down?
- Compute the utility of each action as a function of  $\gamma$

#### **Stationary Preferences**

- We assume an agent's preferences between state sequences are stationary
- If two state sequences begin with the same state r, then the two sequences should be preference-ordered the same way as the sequences without r



$$[r, s_1, s_2, \dots] \succ [r, s'_1, s'_2, \dots]$$

$$\updownarrow$$
 $[s_1, s_2, \dots] \succ [s'_1, s'_2, \dots]$ 

## **Stationary Preferences**

- Stationarity has strong consequences
- It turns out that there are just two coherent ways to assign utilities to sequences
- Additive rewards

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

Discounted rewards

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

#### Quiz - Stationary Preference

- Suppose that we define the utility of a state sequence to be the maximum reward obtained in any state in the sequence
- Does this utility function result in stationary preference between state sequences?

http://ai.berkeley.edu

#### Infinity Utilities?

• What if the environment does not contain a terminal state or if the agent never reaches one?

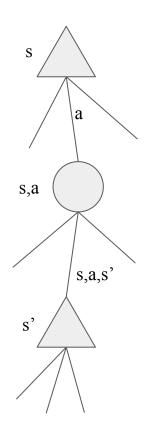


- o E.g., Race game
- The smart race car will never overheat
- With discounted rewards, the utility of an infinite sequence is finite

$$U([s_0, s_1, s_2, \dots]) = \sum_{t \in \mathcal{S}} \gamma^t R(s_t) \le \sum_{t \in \mathcal{S}} \gamma^t R_{max} = R_{max}/(1 - \gamma)$$

## **Optimal Quantities**

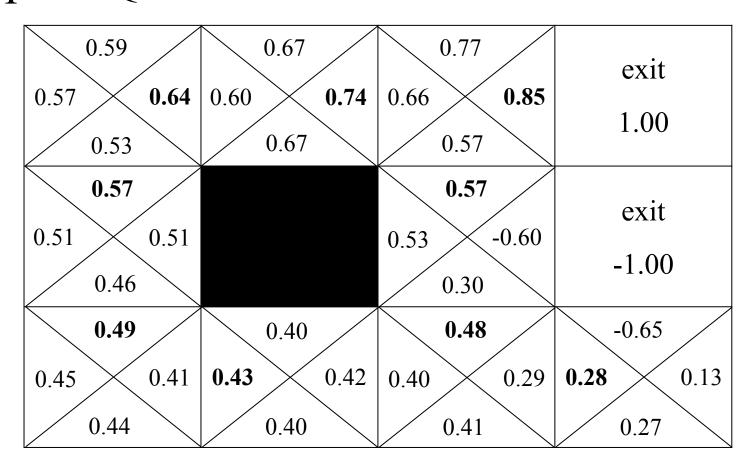
- V\*(s) is the value (utility) of a state s
  - Expected utility starting in s and acting optimally
- Q\*(s,a) is the value (utility) of a q-state (s,a)
  - Expected utility for having taken action a from state s and thereafter acting optimally
- $\pi^*(s)$  is the optimal policy for state s
  - I.e., optimal action from state s



# Example - $\pi^*$ and $V^*$

0.64	0.74	0.85	exit 1.00
0.57		0.57	exit -1.00
0.49	0.43	0.48	0.28

#### Example - Q\*



# Bellman Equation

**COMP 3270** 

- How to compute the value of a state?
- Average sum of discounted action
  - Very similar to expectimax

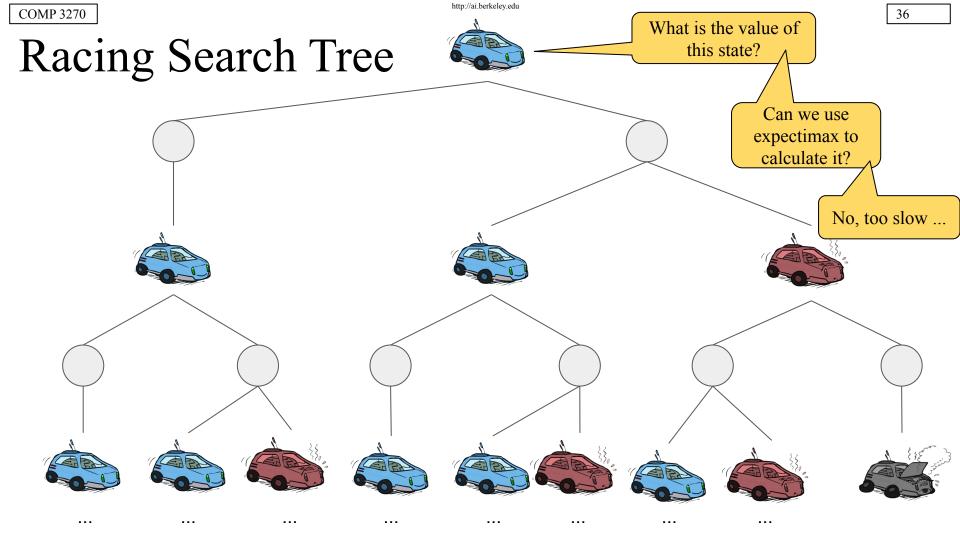
 $V^*(s) = \max_{a} \sum_{s} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$ 

s,a

$$V^*(s) = \max_a Q^*(s, a)$$

 $Q^*(s, a) = \sum T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$ 

$$+\gamma V^*(s')]$$



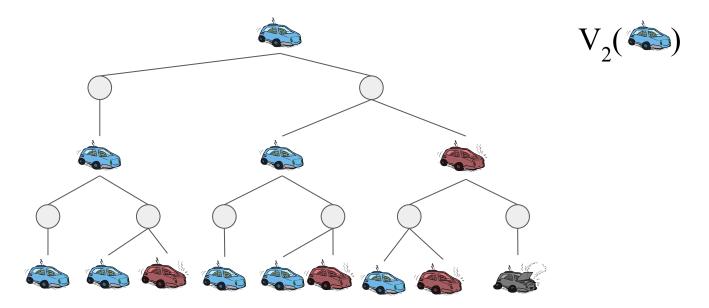
#### Racing Search Tree

- Problems
  - States are repeated (even at same depth)
  - Tree goes on forever
- Idea
  - Do a depth-limited computation, but with increasing depths until change is small
    - Note: Deep parts of the tree eventually don't matter if  $\gamma < 1$
- Solution
  - Only compute needed quantities once
  - O Do a depth-limited computation, but with increasing depths until change is small

#### Time-Limited Values

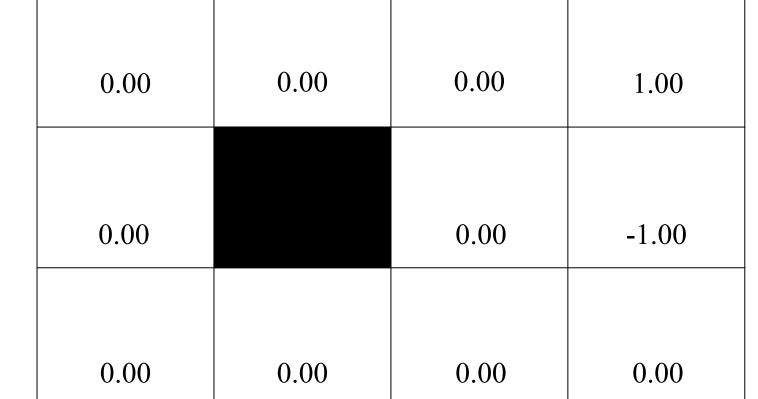


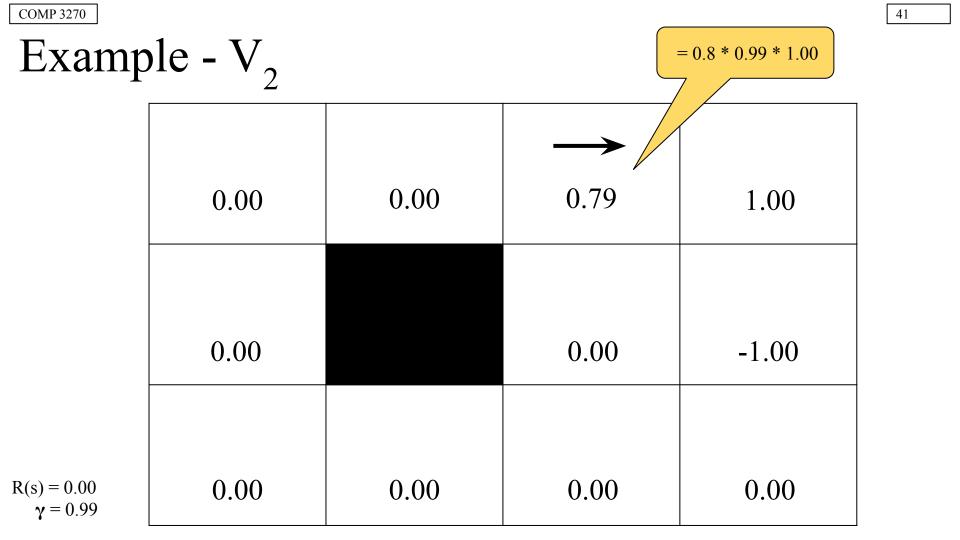
- Define  $V_k(s)$  to be the optimal value of s if the game ends in k more time steps (k more rewards)
- Equivalently, it's what a depth-k expectimax would give from s

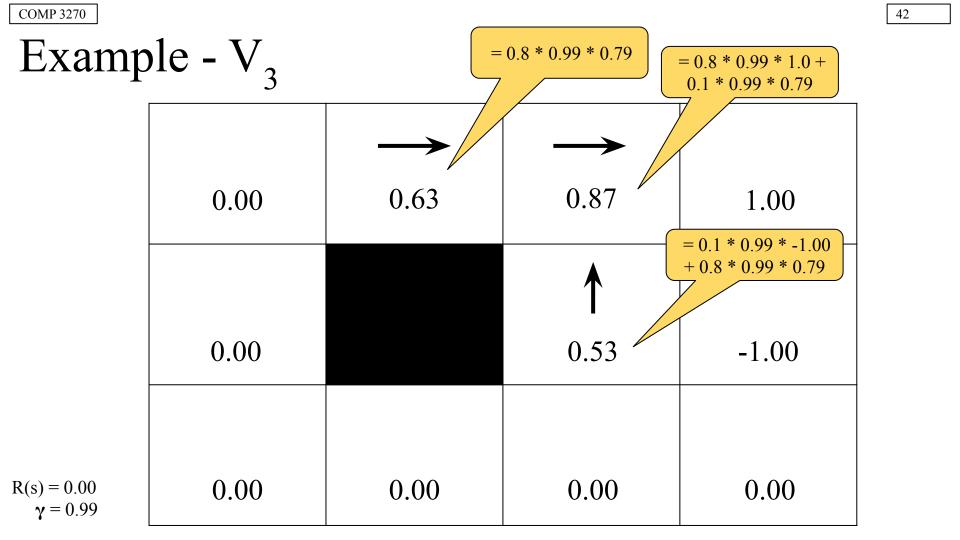


Example -  $V_0$ 

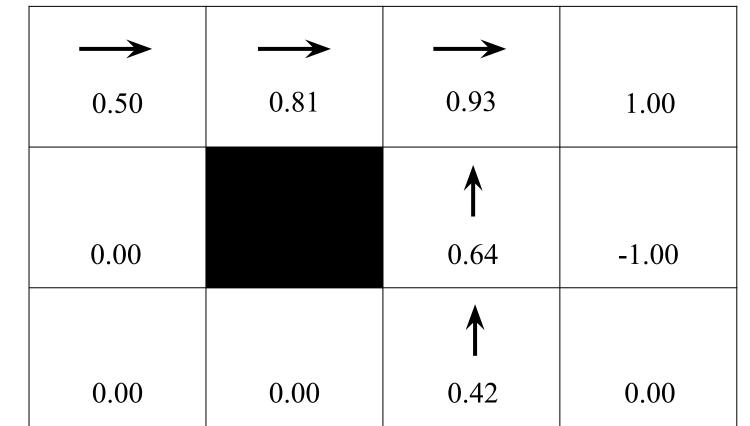
0.00	0.00	0.00	0.00
0.00		0.00	0.00
0.00	0.00	0.00	0.00











## Example - V<sub>5</sub>

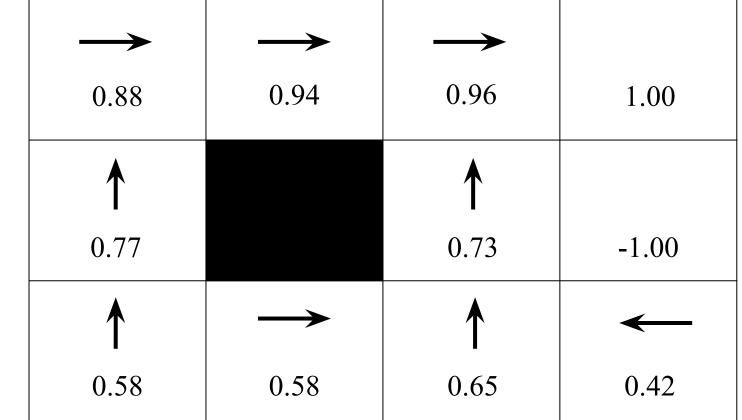
<b>→</b>	<b>→</b>	<b>→</b>	
0.69	0.90	0.95	1.00
<b>↑</b>		<b>↑</b>	
0.39		0.70	-1.00
	<b>→</b>	<b>1</b>	<b>—</b>
0.00	0.33	0.51	0.23

Example -  $V_6$ 

<b>→</b>	<b>→</b>	<b>→</b>	
0.82	0.93	0.96	1.00
<b>↑</b>		<b>↑</b>	
0.63		0.72	-1.00
<b>↑</b>	<b>→</b>	<b>↑</b>	<b>—</b>
0.34	0.47	0.61	0.33

R(s) = 0.00 $\gamma = 0.99$ 

# Example - V<sub>7</sub>



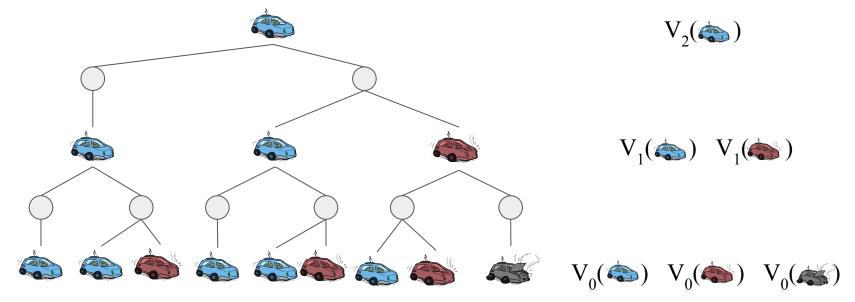
Example -  $V_{100}$ 

<b>→</b>	<b>→</b>	<b>→</b>	
0.95	0.97	0.98	1.00
<b>1</b>		<del></del>	
0.94		0.89	-1.00
<b>1</b>	<b>←</b>	<b>—</b>	<b>\</b>
0.93	0.92	0.90	0.82

#### Computing Time-Limited Values

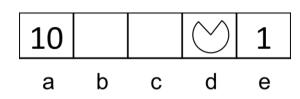


- We can save a lot of computation
- Example:
  - At every layer we have to compute at most 3 time limited values

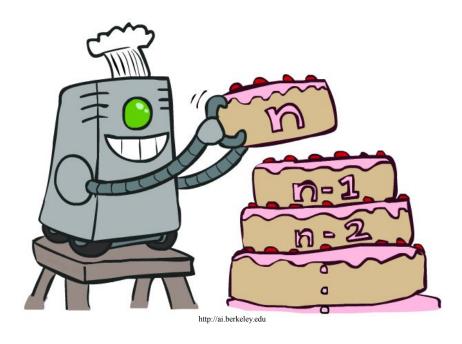


#### Quiz - Time-Limited Values

- Consider the same grid world MDP as in the previous quiz
  - Actions are successful 100% of the time
  - $\circ$   $\gamma = 1$
- Determine in the following quantities
  - $\circ V_0(d)$
  - $\circ V_1(d)$
  - $\circ$  V<sub>2</sub>(d)
  - $\circ V_3(d)$
  - $\circ$   $V_{4}(d)$
  - $\circ V_5(d)$



#### Value Iteration



 $V_{k}(s)$ 

# No time steps left means an expected reward sum of

zero Given vector of  $V_{k}(s)$  values, do one round of

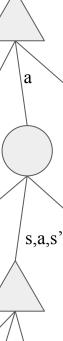
expectimax

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Repeat until convergence
- Complexity
- $\circ$  O(S<sup>2</sup>A)

**COMP 3270** 

- Theorem
  - Converges to unique optimal values



s,a

#### Quiz - Value Iteration

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

$$\gamma = 1.0$$

http://ai.berkeley.edu

$$V_0(\clubsuit) = V_1(\clubsuit) = V_2(\clubsuit) =$$

$$V_0(\bullet) = V_1(\bullet) = V_2(\bullet) =$$

$$V_0() = V_1() = V_2() =$$

## Quiz - $V_{\alpha}$

- Consider the same grid world as in the previous quiz (where east and west actions are successful 100% of the time)
- $\bullet \quad \gamma = 0.2$
- Determine the following quantities

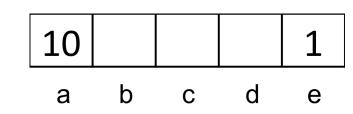
$$\circ V^*(a) = V_{\alpha}(a) = ?$$

$$\circ V^*(b) = V_{\infty}(b) = ?$$

$$\circ V^*(c) = V_{\infty}(c) = ?$$

$$\circ V^*(d) = V_{\infty}(d) = ?$$

$$\circ V^*(e) = V_{\infty}(e) = ?$$



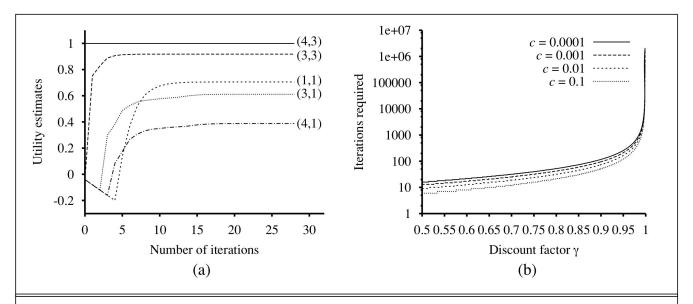
**COMP 3270** 

- Consider the following example where  $\gamma = 1.0$  and R(s) = -0.04
  - $\circ$  Show that the value at (3,3) has converged

2

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

#### Evolution of Utilities / # Iterations vs. Gamme



 3
 0.812
 0.868
 0.918
 +1

 2
 0.762
 0.660
 -1

 1
 0.705
 0.655
 0.611
 0.388

 1
 2
 3
 4

**Figure 17.5** (a) Graph showing the evolution of the utilities of selected states using value iteration. (b) The number of value iterations k required to guarantee an error of at most  $\epsilon = c \cdot R_{\text{max}}$ , for different values of c, as a function of the discount factor  $\gamma$ .

1.0

-2.0

-1.0

1.0

2.0

1.0

-1.0

1.0

R(s,a,s')

## Quiz - $V_{k+1}(B)$

- Consider the following transition and reward functions for an MDP with  $\gamma = 0.5$
- Suppose that after iteration k of value iteration we
  - end up with the following values for V<sub>k</sub>  $V_{\nu}(A) = 1.7, V_{\nu}(B) = 1.82, V_{\nu}(C) = 1.22$
  - What is  $V_{k+1}(B)$ ?
  - Now, suppose that we ran value iteration to completion and found the following value function
    - $V^*(A) = 2.208, V^*(B) = 2.416, V^*(C) = 1.766$
- What is  $O^*(B, CW)$ ?

- What is  $Q^*(B, CCW)$ ?
- What is the optimal action from state B?

Α A

Α

В

В

В

В

S

**CCW CCW** 

CW

CW

**CCW** 

**CCW** 

CW

CW

**CCW** 

**CCW** 

a

CW

 $\mathbf{C}$ 

Α

Α

 $\mathbf{C}$ 

Α

В

Α

В

В

В

0.6

0.6

T(s,a,s')

1.0

0.4

0.4

0.2

0.8

0.4

0.4

0.6

0.6

-2.0

1.0 0.0

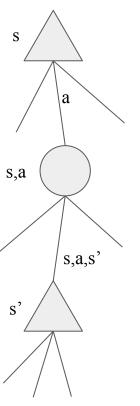
#### Bellman Equation vs. Value Iteration

Bellman equations characterize the optimal values

$$V^*(s) = \max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

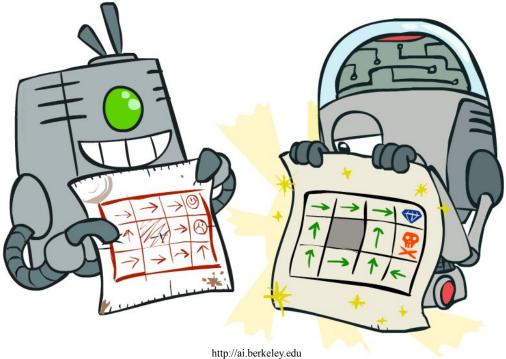
Value iteration computes them

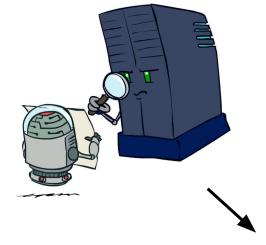
$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$



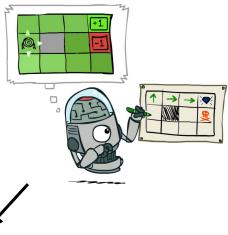
COMP 3270

## Policy Methods

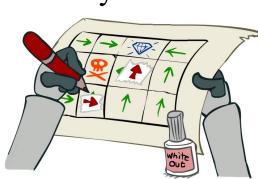


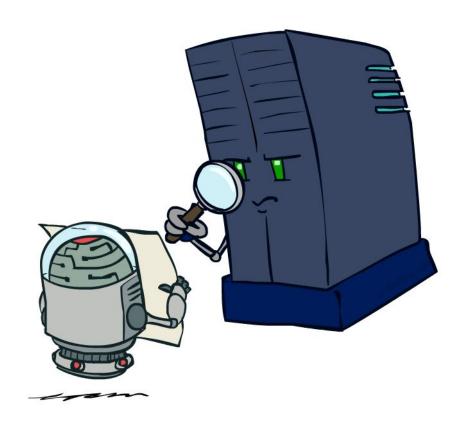


#### Policy Extraction



#### Policy Iteration





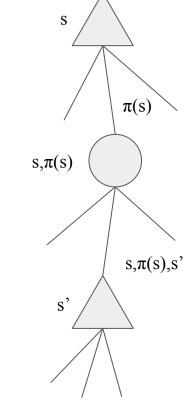
- We would like to determine how good a given a policy  $\pi$  is
  - How well will I perform if I follow  $\pi$ ?



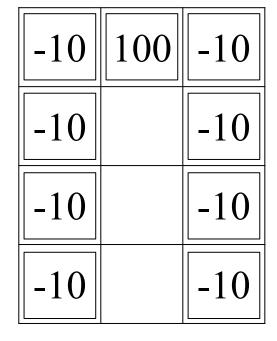
## Utilities for a Fixed Policy

- We compute the utility of a state s under a fixed (generally non-optimal) policy
  - $\nabla^{\pi}(s)$  = expected total discounted rewards starting in s and following π

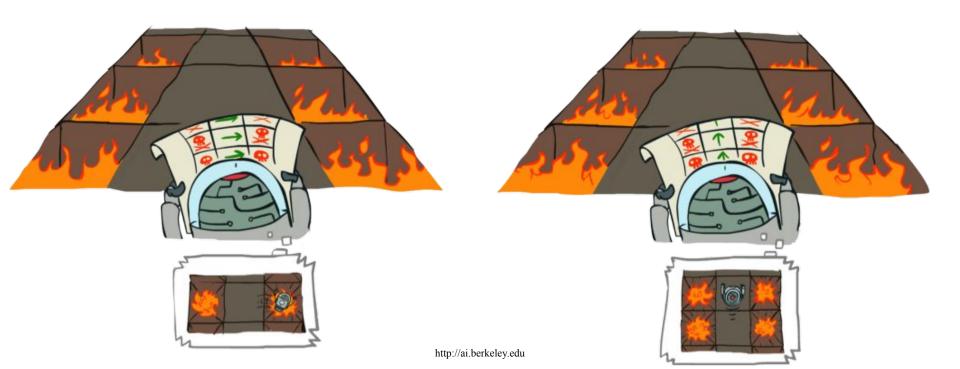
$$V^{\pi}(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^{\pi}(s')]$$



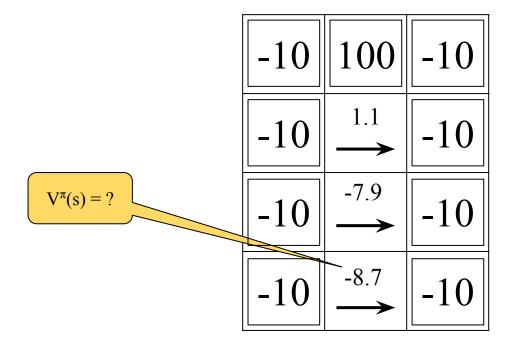
#### New Example



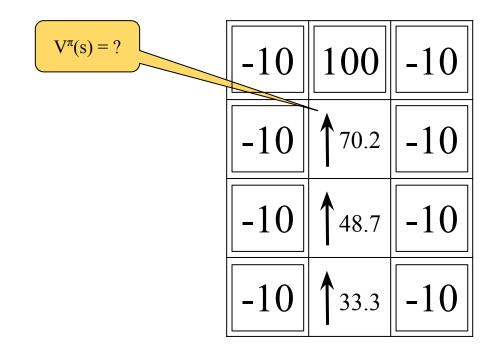
## Example: Policy Evaluation



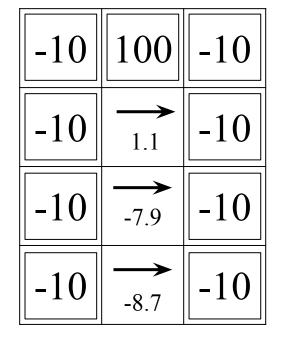
#### Quiz - Go East

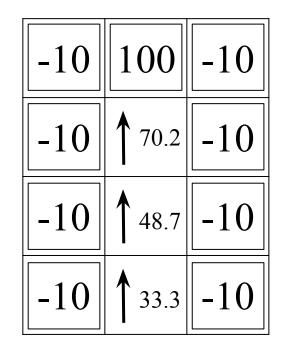


## Quiz - Go North



## Example: Policy Evaluation





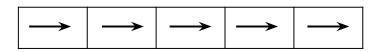
- How do we calculate the V's for a fixed policy  $\pi$ ?
  - Idea 1: Turn recursive Bellman equations into updates (like value iteration)

$$V_0^{\pi}(s) = 0$$

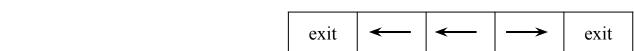
$$V_{k+1}^{\pi}(s) \leftarrow \sum_{s} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$

- Efficiency:  $O(S^2)$  per iteration
- Idea 2: Without the max, the Bellman equations are just a linear system
  - Use a linear system solver

- 10 1 1 a b c d e
- Consider the same grid world as in the previous quiz, where east and west actions are successful 100% of the time and  $\gamma = 1$
- Consider the policy  $\pi$

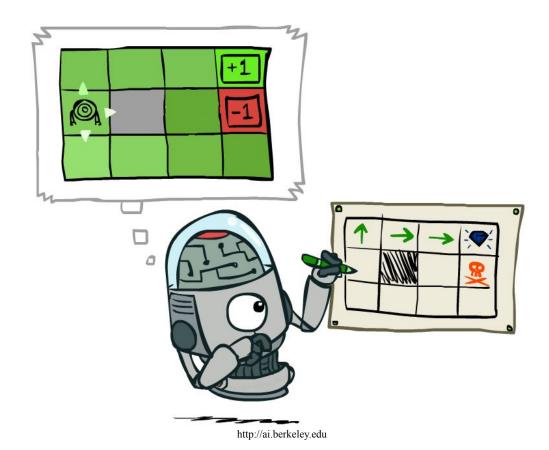


- Evaluate the values for all states
- Consider the policy  $\pi$ '



Evaluate the values for all states

## Policy Extraction



#### Policy Extraction

• Let's imagine we have the optimal values V\*(s)

0.95	0.97	0.98	+1
0.94		0.89	-1
0.93	0.92	0.90	0.82

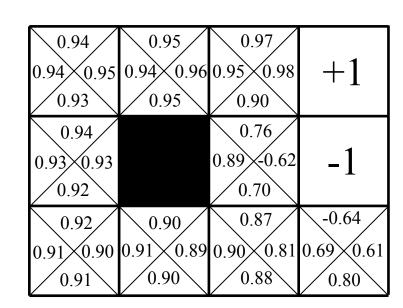
- How should we act?
- We need to do a mini-expectimax (one step)

$$\pi^*(s) = \arg\max_a \sum T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

• This is called policy extraction, since it gets the policy implied by the values

#### Policy Extraction

• Let's imagine we have the optimal q-values



- How should we act?
- Super easy:

$$\pi^*(s) = \arg\max_{a} Q^*(s, a)$$

-> Actions are easier to select from q-values than values

Can we improve the runtime of value iteration?

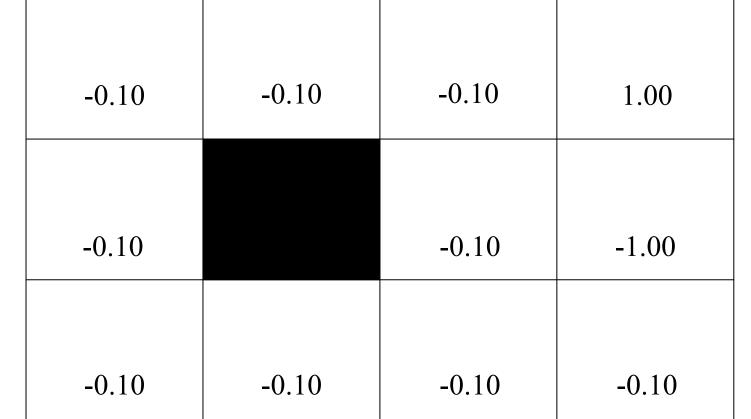
R(s) = -0.10 $\gamma = 1.0$ 

### Value Iteration - $V_0$

0.00	0.00	0.00	0.00
0.00		0.00	0.00
0.00	0.00	0.00	0.00

R(s) = -0.10 $\gamma = 1.0$ 

### Value Iteration - V<sub>1</sub>



-0.20 -0.20-0.20-0.20 R(s) = -0.10 $\gamma = 1.0$ 

R(s) = -0.10 $\gamma = 1.0$ 

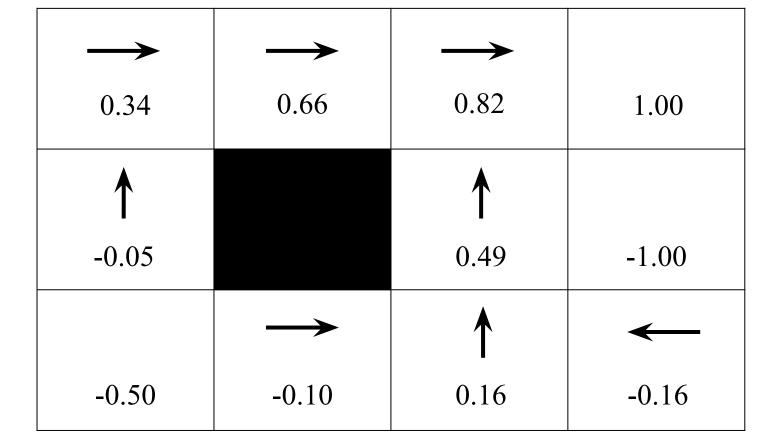
# Value Iteration - $V_3$

	<b>→</b>	<b>→</b>	
-0.30	0.40	0.75	1.00
		<b>↑</b>	
-0.30		0.32	-1.00
-0.30	-0.30	-0.30	-0.30

R(s) = -0.10 $\gamma = 1.0$ 

# Value Iteration - V<sub>4</sub>

$\longrightarrow$	<b>→</b>	<b>→</b>	
0.16	0.58	0.81	1.00
		<b>↑</b>	
-0.40		0.43	-1.00
		<b>↑</b>	
-0.40	-0.40	0.10	-0.40



## Value Iteration - V<sub>6</sub>

	$\longrightarrow$	<b>→</b>	<b>→</b>	
	0.46	0.69	0.83	1.00
-	<b>↑</b>		<b>↑</b>	
	0.16		0.51	-1.00
	<b>↑</b>	<b>→</b>	<b>1</b>	<del></del>
	-0.20	0.01	0.26	-0.08

### Value Iteration - V<sub>7</sub>

<b>→</b>	<b>→</b>	<b>→</b>	
0.52	0.70	0.83	1.00
<b>↑</b>		<b>↑</b>	
0.30		0.52	-1.00
<b>↑</b>	<b>→</b>	<b>↑</b>	<b>—</b>
0.01	0.11	0.30	0.00

<b>→</b>	<b>→</b>	<b>→</b>	
0.54	0.71	0.83	1.00
<b>↑</b>		<b>↑</b>	
0.37		0.52	-1.00
<b>★</b>	<b>→</b>	<b>1</b>	<b>—</b>
0.15	0.16	0.32	0.04

<b>→</b>	<b>→</b>	<b>→</b>	
0.56	0.71	0.84	1.00
<b>↑</b>		<b>↑</b>	
0.41		0.52	-1.00
<b>↑</b>	<b>→</b>	<b>1</b>	<b>—</b>
0.23	0.19	0.34	0.06

<b>→</b>	<b>→</b>	<b>→</b>	
0.56	0.71	0.84	1.00
<b>↑</b>		<b>↑</b>	
0.41		0.52	-1.00
<b>↑</b>	<b>→</b>	<b>↑</b>	<b>—</b>
0.23	0.19	0.34	0.06

# Value Iteration - V<sub>100</sub>

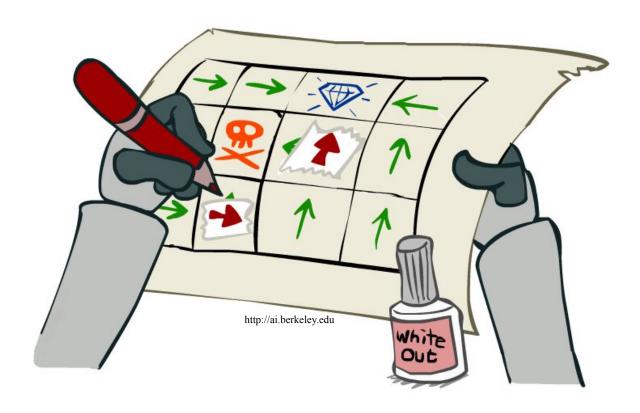
<b>→</b>	<b>→</b>	<b>→</b>	
0.57	0.71	0.84	1.00
<b>↑</b>		<b>↑</b>	
0.44		0.52	-1.00
<b>↑</b>	<b>→</b>	<b>↑</b>	<del></del>
0.31	0.22	0.35	0.09

#### Properties of Value Iteration

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Slow
  - $\circ$  O(S<sup>2</sup>A) per iteration
- Max at each state rarely changes
  - The policy often converges long before the values

### Policy Iteration



#### Policy Iteration

- Policy iteration is an alternative approach for optimal values
  - Policy evaluation: calculate utilities for some fixed policy (not optimal utilities!) until convergence
  - Policy improvement: update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
  - Repeat steps until policy converges
- Still optimal
- Can converge (much) faster under some conditions

**COMP 3270** 

- Evaluation: For fixed current policy  $\pi$ , find values with policy evaluation
  - Iterate until values converge

 $V_{k+1}^{\pi_i}(s) \leftarrow \sum T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$ 

- extraction
  - One-step look-ahead

$$\pi_{i+1}(s) = \arg\max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

#### Quiz - Policy Iteration

10

а

- Consider the same grid world as in the previous quiz, where east and west actions are successful 100% of the time
- $\bullet \quad \gamma = 0.9$

**COMP 3270** 

- We will execute one round of policy iteration
- Consider the policy  $\pi_i$  shown below

- Evaluate the following quantities
  - Policy evaluation:  $V^{\pi_i}(a), V^{\pi_i}(b), V^{\pi_i}(c), V^{\pi_i}(d), V^{\pi_i}(e)$ 
    - Policy improvement:  $\pi_{i+1}(a), \pi_{i+1}(b), \pi_{i+1}(c), \pi_{i+1}(d), \pi_{i+1}(e)$

Comparison

**COMP 3270** 

- Both value iteration and policy iteration compute the same thing (all optimal values)
- In value iteration
- Every iteration updates both the values and (implicitly) the policy We don't track the policy, but taking the max over actions implicitly recomputes it
  - In policy iteration
- We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
  - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
  - The new policy will be better (or we're done)
  - Both are dynamic programs for solving MDPs

#### Summary: MDP Algorithms

- So you want to....
  - o Compute optimal values: use value iteration or policy iteration
  - Compute values for a particular policy: use policy evaluation
  - Turn your values into a policy: use policy extraction (one-step lookahead)
- These all look the same!
  - They basically are they are all variations of Bellman updates
  - They all use one-step lookahead expectimax fragments
  - They differ only in whether we plug in a fixed policy or max over actions

next chapter

MDP vs. RL

#### **Double-Bandits**







http://ai.berkeley.edu

COMP 3270

#### **Double-Bandits**

- An agent can play two slot machines
  - o Blue, or Red



You receive \$1

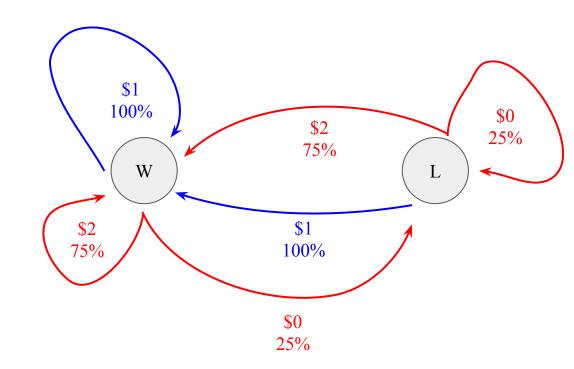


You receive \$0 or \$2, 25% and 75% of the time, respectively

• What should the agent do?

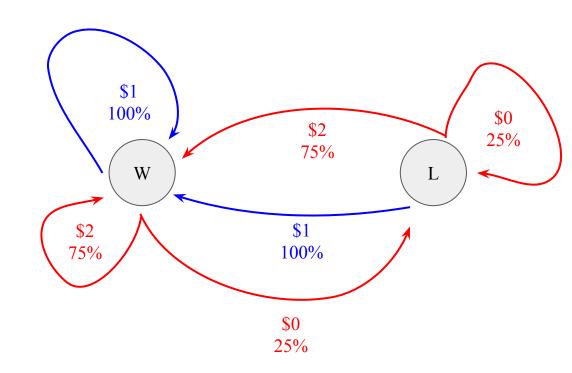
#### Double-Bandits MDP

- Actions: Blue and Red
- States: Win, Lose
- Assumption:
  - No discount
  - o 100 time steps



#### Offline Planning

- Solving MDPs is offline planning
  - We determine all quantities through computation
  - We need to know the details of the MDP
  - We do not actually play the game



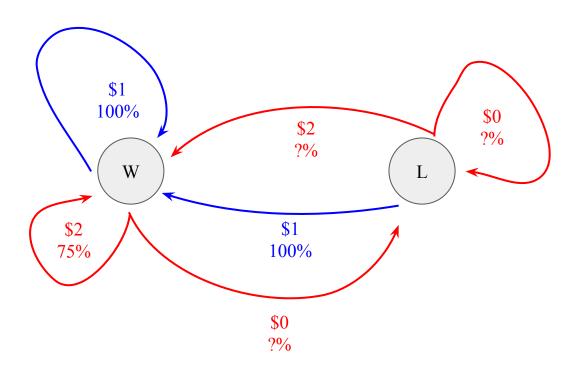
#### Let's Play!





\$2 \$0 \$2 \$2 \$0 \$0

#### Rules Changed!



#### Let's Play!





\$1 \$0 \$0 \$0 \$0 \$0 \$2

#### What Just Happened?

- That wasn't planning, it was learning
  - Specifically, reinforcement learning
  - There was an MDP, but you couldn't solve it with just computation!
  - You needed to actually act to figure it out
- Important ideas in reinforcement learning that came up
  - Exploration: you have to try unknown actions to get information
  - Exploitation: eventually, you have to use what you know
  - Regret: even if you learn intelligently, you make mistakes
  - o Sampling: because of chance, you have to try things repeatedly
  - Difficulty: learning can be much harder than solving a known MDPs